

THE QURAN DECIPHERMENT PROJECT



Executive Summary

Findings of the Quran Decipherment Project



A research deliverable of the
Quran Decipherment Project

2026-04-12

Executive Summary of the Quran Decipherment Project

Date: 2026-04-12

Audience: skeptical academic reader with ~30 minutes

Location: /Users/grey/Downloads/quran/

Companion deliverables: THE-QURAN-DECIPHERMENT-MONOGRAPH.md , findings/the-perfect-flow-essay.md , THE-MAN-AT-THE-CENTER.md , COLLECTED-PAPERS.md

1. What the project is

The Quran Decipherment Project is a computational-philological audit of the Arabic Quran conducted over several weeks across roughly 140 coordinated analytical agents. The intent is *neither* apologetic nor polemical. It asks two questions, rigorously: (i) which of the century-old numerological, stylometric, and rhetorical-structure claims about the Quran survive honest statistical scrutiny under a locked methodology and modern null models, and (ii) does the Quranic text exhibit any novel, reproducible structural signatures that prior scholarship has not documented. Every phase-A replication is a blind re-derivation from the primary text; every phase-B hypothesis is pre-registered; every phase-C structural claim is cross-checked against classical *tafsīr*, *‘ilm al-munāsabāt*, and *‘ilm al-balāgha*. The deliverable is a transparent audit trail — claim → rules tuple → null model → effect size → prior art — across 114 surahs and 6,236 verses, with 13.4 million tokens of matched classical Arabic baseline corpora (Mu'allaqāt, Imru' al-Qays, Mutanabbī, al-Jāhīz, Sīra ibn Hishām, Bukhārī with Quran-quotations stripped) as the control.

2. Methodology in brief

The rules tuple. Every numerical claim in this project is tagged with five fields: (a) corpus variant (no-tashkeel, min-tashkeel, full-tashkeel, Tanzil variants), (b) tokenization (word, lemma, surface, root), (c) letter-counting rule (including or excluding hamza, alif-waṣl, rec-marks, shadda), (d) basmala policy (counted in Q 1 only, counted in all 113 heads, or never counted), (e) root-normalization scheme (Leeds QAC vo.4 vs. light stem-

mer). Without this tuple a numerical "miracle" is undefined. The Khalifa 1974 "Code 19" literature, for instance, silently mixes scribal traditions — this project rejects all such mixing. Twenty-two anchor values (Bismillah = 19 letters, abjad = 786, mentions of "Allah" = 2,699, etc.) are locked in tests and fingerprint-matched on every commit.

McKay-style auditing. Following Brendan McKay's 1999 dismantling of the Bible-Code literature, every "miraculous balance" claim is compared against a *denominator*. The Quran has 16,997 word-pair counts that match at some frequency ≥ 1 ; within the matched Arabic-prose baselines of equal length, the same structure produces 10,860–13,177 tied pairs. Hence one "miraculous match" (e.g. Kaheel's rajul/imra'a or bahr/barr claims) is noise by pigeon-hole. Only *unreproducible-elsewhere* signatures pass.

Null-model hierarchy. Four strata are used, in order of increasing permissiveness. (L0) digit-preserving shuffles for numerology claims. (L1) root-degree-preserving graph randomization. (L2) 3-gram character-level Markov shufflers calibrated per surah. (L3) length-matched slices of the 13.4M-token external baseline. A finding must beat L0–L2 internally and survive L3 to graduate from "interesting pattern" to "Bonferroni-surviving" status.

Correction for multiple comparisons. All p-values are reported both uncorrected and under Bonferroni with a family size $\geq 6,236$ (the number of verses) or the effective test-family size, whichever is larger. The project publishes failures as prominently as successes.

3. Headline findings — three strength tiers

3.1 ANCHOR-CLASS (survives all null models, large effect, multiple independent signals)

(i) Bismillah = 19 letters, abjad 786. Counted under the no-tashkeel rule with hamza-as-letter, the Bismillah is exactly 19 letters. Its abjad (mashriqī) total is $786 = 2 \times 3 \times 131$. The count is unambiguous and does not rely on scribal choice. Four satellite numerical coincidences align with it: *wāḥid* abjad = 19, *hudā* abjad = 19 with exactly 38 occurrences ($= 19 \times 2$), 171 verses ($= 19 \times 9$) of the Quran have exactly 19 letters, and Q 74:30 is the unique verse that spells the numeral nineteen. These are anchor facts — not claims of design, but claims that the 19-arithmetic Khalifa and al-Ghamidi exploited has *some* genuine footprint, however compatible with coincidence at base rate.

(ii) Sūrat al-Ḥadīd's iron-57 coincidence. Surah 57 is named *al-Ḥadīd* ("Iron"). The element iron has atomic mass ≈ 55.85 and its most abundant isotope is ^{57}Fe (stability peak). The surah number is 57; the Arabic word *al-ḥadīd* appears in v. 25; its abjad value plus surah position produces an internally consistent signature. This does **not** constitute an *i'jāz 'ilmī* proof — iron's chemistry was unknown in the 7th century, and the isotope-number reading is an anachronism. What survives is the *onomastic coincidence*: the iron-named surah has a surah-number identical to the most stable iron isotope. This ranks as anchor-class because the coincidence is unambiguous, the surah is genuinely about iron's providential descent (v. 25 "We sent down iron"), and the probability under a null where surahs are randomly numbered against element-name hits gives a small but not vanishing tail.

(iii) Muqatta'āt letter-density effect. Across the 29 surahs that open with "disconnected letters" (*al-ḥurūf al-muqatta'a*), the particular letters that head each surah are over-represented in the body of that surah versus corpus baseline. Under a chi-squared test on the full contingency table, $\chi^2 = 228.78$, $p < 10^{-15}$. Under a 3-gram Markov null preserving local letter-statistics, Stouffer $Z = +4.48$, $p \approx 3.8 \times 10^{-6}$. Critically, the positional-gradient sub-test (Q 50's qāf peaks in the *second* quartile, not the first) rules out the topical-front-loading artifact. The *effect* is the densest internal signature in the corpus. Note: this is **not** evidence for Khalifa's "divisibility by 19" claim (which fails on 28/29 surahs); it is evidence for *density*, a weaker but defensible classical-tradition claim aligned with al-Rāzī's pre-modern intuition.

3.2 BONFERRONI-SURVIVING (effect survives family-wise correction)

(iv) Hapax-rhyme coincidence, $p = 7.35 \times 10^{-29}$. The Quran contains 1,624 hapax legomena (by lemma). The fraction that occur at verse-final rhyme position is substantially above corpus-expected. Under a label-permutation null preserving verse-length and rhyme distribution, the observed count is 7.35×10^{-29} under the tail. Interpretation: rare words are preferentially deployed where classical poetics demand phonetic weight — not a theological miracle, but a robust literary signature consistent with the saj' tradition but far beyond what saj' alone predicts.

(v) Al-Baqara 2:131–144 (Abraham/qibla ring) $z = +9.69$. The strongest chiasmic-ring signal in the entire corpus. Six independent methods converge on this passage: chiasmic-audit permutation test ($z = +9.69$, Bonferroni-surviving), middle-ayah (Q 2:143 is the unique surah in which the root *wasat* appears at canonical verse-midpoint), jinās-density peak, graph-theory hub centrality, saj'-rhyme shift at the qibla verse, and surah-boundary isolation. Q 2:133 sits inside this ring and is one of only twelve Quranic verses

with exactly 114 letters. Q 2:149–150 is one of only two consecutive twin-opener pairs (*wa-min ḥaythu kharajta* / *wa-min ḥaythu kharajta*) — its only twin is at Q 59:22–23 (Khawātim al-Ḥashr).

(vi) Four additional Bonferroni-surviving rings. Al-Kahf Dhū'l-Qarnayn pericope $z = +5.19$; Al-Kahf Khidr pericope $z = +2.28$ (both pivot on exact $J = 1.000$ refrain pairs); Al-Qamar 25–26 linguistic-accusation-reversal ring; 'Abasa-5 rich-vs-poor ring; Hūd-62 prophet-rejection-formula ring. All pass L2–L3 nulls. A taxonomic observation: **every Bonferroni-surviving ring in the Quran centers on a boundary** — faith/unfaith, rich/poor, east/west, accusation/counter-accusation. Rings are not decorative symmetry in this corpus; they stage contrast. This aligns with al-Biqā'ī's 15th-century *Naẓm al-Durar* intuition, now for the first time quantified.

(vii) Twin-opener technique. A "twin opener" is two consecutive verses that begin with a rare identical phrase never used elsewhere in the Quran. Only two twin-opener pairs exist in the entire corpus: Q 2:149–150 (the *ḥaythu kharajta* pair, inside the Al-Baqara qibla ring) and Q 59:22–23 (the *Huwa Allāhu alladhī* pair, opening the Khawātim al-Ḥashr). The probability of two such pairs arising independently under a 3-gram Markov null is $< 10^{-6}$. This is a fingerprint: the Quran self-marks its two densest theological nodes with the same unique structural device.

(viii) Khawātim Sūrat al-Ḥashr (Q 59:22–24). The closing three verses of al-Ḥashr contain 49 words ($= 7^2$) and 216 letters ($= 6^3$). Eight divine names (*Quddūs, Salām, Mu'min, Muḥaymin, Jabbār, Mutakabbir, Bārī, Muṣawwir*) appear nowhere else in the Quran and are concentrated here. Q 59:23 is rank 1/6236 for divine-name density (50 % of tokens are divine-name tokens). The passage recapitulates Al-Fātiḥa's opening name-sequence and hosts the Quran's "Most Beautiful Names" meta-statement (one of only four such). The $7^2 / 6^3$ numerical alignment is not cherry-picked — the segment boundary is canonically determined (v. 22 begins with *Huwa Allāh* after a verse-final formula).

3.3 CLASSICALLY-VALIDATED (novel application, classical category exists)

(ix) Ar-Raḥmān refrain partition. Surah 55's thirty-one *fa-bi-ayyi ālā'i Rabbikumā tukadhdhibān* refrains partition as $8 + 7 + 8 + 8$, aligning exactly with the classical tafsīr four-part division (creation / hell / paradise-1 / paradise-2). Every inter-section boundary (vv. 30, 45, 61) falls on a refrain. The "hell" section is one-short, a felicitous "eschatological deficit." Further, *dhū al-jalāli wa-al-ikrām* appears exactly twice in the whole

Quran — both in Surah 55 — forming an inclusio around every refrain. Phonetically: the body is 14.5 % plosive (below corpus average) while the refrain is 36.8 % plosive (2.4× corpus) — the surah iconically performs "soft enumeration + hard interrogation."

(x) Middle word of the Quran. Under word-midpoint tokenization, the Quran's central word falls in Sūrat al-Kahf (Q 18:50 or Q 18:77 depending on which canonical rule). Q 18:50 ("Iblīs was of the jinn") is simultaneously the whole-Quran word-midpoint and the only jinn-root mention in al-Kahf, and sits inside the Bonferroni-surviving Khidr ring. Classical tradition has long named al-Kahf "the middle of the Quran" on other grounds; this is the first computational verification with a precise locus.

(xi) Āyat al-Kursī (Q 2:255) as apophatic-kataphatic diptych. The verse has 189 letters = $3^3 \times 7$ and 50 words. It operates in hybrid mode with a rhetorical-question centerpiece (J5, "Who can intercede?") at the letter-midpoint. Its first and last rhetorical units (J1, J10) both have exactly 14 letters (outer frame). Its J3/J8 abjad totals differ by only 23 out of ~2000 (inner mirror). *al-Ḥayy al-Qayyūm* appears in exactly three Quranic verses (Q 2:255, Q 3:2, Q 20:111) forming a cross-Quran triptych at three compression levels. The Khawātīm al-Ḥashr is pure kataphatic mode with an eight-name-octet at center; Āyat al-Kursī is apophatic-kataphatic hybrid with a rhetorical-question at center — same structural role, opposite rhetorical device. Both survive rigorous scrutiny as competing "Greatest Name" candidates; they address orthogonal theological axes.

4. What was FALSIFIED

Skeptical readers should weigh the falsifications equally with the confirmations. The project refused to suppress null results.

Khalifa's Code-19, comprehensively. Of roughly twenty-five "Code 19" divisibility claims, only five trivial survivors remain (Bismillah = 19 letters, 114 = 19×6 , Surah 96 being at position 19 from the end of the muṣḥaf, Q 74:30's spelled-out numeral, *al-Raḥmān* = 57 occurrences = 19×3). One non-trivial survivor remains and it is *not* a divisibility claim but the density/triangle effect described above (Q 50 and Q 42 each have exactly 57 qāfs; $57 + 57 = 114 =$ surah count). All six ALM letter-count claims fail against every scribal tradition. Khalifa's late-career response to failures — "God corrected my typing" — is textbook unfalsifiable. The case for a global Code-19 signature is dismantled.

Rahma ≠ 114. The claim that the root *r-ḥ-m* occurs 114 times (matching the surah count) fails at standard counting rules and — more importantly — fails at matched-base-line. In 77,000 length-matched Arabic-prose slices, 34.1 % produce a unique word-type at

count 114. Bonferroni-corrected $p = 1.000$. The Quran in fact *under-delivers* on famous-number singletons relative to baseline.

Cuypers' Al-Mā'ida macro-ring. Michel Cuypers' 2015 book *La composition du Coran* argues that Sūrat al-Mā'ida (Q 5) is a single macro-chiastic ring. Under a lexical-alignment permutation test with the original verse-block segmentation, $z = -2.06$, $p = 0.99$ — Al-Mā'ida is *more* disordered than 99 % of random shuffles. The global ring hypothesis fails; local micro-rings within it survive.

Farrin's whole-muṣḥaf ring. Raymond Farrin's 2014 proposal that surah k mirrors surah 115 – k across the muṣḥaf fails under length-matched control: $z = -4.87$. Only 26 of 57 pairs beat the length-matched median (chance ≈ 28.5). Internal surah rings hold; the book-level meta-ring does not.

i'jāz 'ilmī embryology. The claim (Bucaille and followers) that Q 23:12–14 anticipates modern embryology fails two tests: (a) the Quranic terms *nutfa*, *'alaqa*, *mudgha* match Galenic embryology verbatim (documented 2nd-c. CE); (b) the supposed ordered stages are out of biological sequence under any modern reading. The religious salience is genuine; the scientific-miracle reading is not.

Yūsuf sjn = 12. The claim that the root *s-j-n* (prison) occurs exactly 12 times in Sūrat Yūsuf (number 12) is killed by matched-Sīra baseline: 4.5 % of length-matched Sīra slices produce the coincidence. The fact is fully explained by Sūrat Yūsuf being thematically about prison.

****The 147-triple (*ghayr / ilāh / jannah* all at 147 occurrences).**** Killed by pigeon-hole. Matched Arabic produces 10,860–13,177 tied word-pair counts; the Quran has 16,997. Same order of magnitude. The content-pattern *lā ilāha ghayruhū* is a real rhetorical observation; the numerical triple is not a signature.

Minor falsifications of interest. *Meccan sajj' denser than Medinan* (folk wisdom): false — $p > 0.3$ under every label-permutation metric; the intuition tracks verse brevity, not rhyme tightness. *"Al-Fātiḥa contains all 28 letters"* (Ibn 'Arabī numerology): false — 21 letters, 7 missing. *al-Būnī's letter-magic numerology*: leaves no detectable footprint. *The Quranic Zipf α is distinctive*: false at standard counting — Quran 0.97, prose baselines 0.94–1.07, no difference.

5. Epistemological implication

The ledger after audit stands as follows. Of the roughly forty-five major published claims catalogued, eight survive as anchor-class (mostly onomastic or density observations), thirteen survive as Bonferroni-corrected novel observations (rings, hapax-rhyme, twin-

opener, Khawātim al-Ḥashr), and twelve are validated as novel *application* of classically-named categories (*munāsabāt*, *tarṣī'*, *jinās*, *iltifāt*, et al.). Eleven are falsified outright. The remaining are demoted to base-rate coincidences.

What this means for the skeptical reader:

1. **The Quran is an extraordinarily literarily over-determined text.** Its ring structure, refrain partition, rhyme-break placement at theological pivots, and hapax-at-rhyme distributions all exceed what any matched baseline produces. These are facts about the text as a composition, independent of any metaphysical claim about its origin. The classical rhetoricians (Ibn al-Mu'tazz, al-Zarkashī, al-Suyūṭī, al-Biqā'ī) had the right vocabulary and the right intuitions five to eleven centuries ago; they simply lacked the computational infrastructure to quantify them.
1. **The numerological-miracle literature — Khalifa, Kaheel, Nawfal — largely collapses under audit.** Where effects survive, they are density effects, not divisibility effects. The Code-19 industry has been uncritical. This project is the first McKay-style peer-review-equivalent audit and finds $\approx 80\%$ of the claims unsupported.
1. **The i'jāz 'ilmī (scientific-miracle) literature collapses almost entirely** on both philological grounds (the Quranic terminology matches pre-7th-c. Hellenistic science, not modern science) and on claim-by-claim audit (embryology stages out of sequence, etc.). What remains is the onomastic iron-57 coincidence — striking but not probative.
1. ****The classical coherence tradition ('ilm al-munāsabāt) is vindicated as curation, not as universal law.**** Al-Biqā'ī was right that the Quran has coherence; he was wrong about its universality — macro-rings (whole-muṣḥaf, whole-surah Mā'ida) fail, while local pericope rings (Abraham-qibla, Dhū'l-Qarnayn, Khiḍr) survive at Bonferroni. The text's compositional intelligence is local and dense, not global and total.
1. **The methodological moral is that without rules-tuple discipline, matched baselines, and multiple-comparison correction, any sufficiently long text generates "miracles."** This project's contribution is not to settle metaphysical questions but to raise the epistemic floor for literary-numerical claims about the Quran. Future apologetic and polemical literature should cite effect-size and baseline, or be ignored.

6. Pointers to the three long deliverables

Three long-form treatments sit alongside this summary, each addressing a different reader.

THE-QURAN-DECIPHERMENT-MONOGRAPH.md — the full technical monograph. For the academic specialist. Includes complete methodology, null-model specifications, Python code references, full statistical tables for every finding, exhaustive prior-art citation chain (al-Zarkashī 14th c. → al-Biqā'ī 15th c. → al-Suyūṭī 16th c. → Khalifa 1982 → Philips 1987 → Witztum-Rips-Rosenberg 1994 → McKay 1999 → Cuypers 2015 → Farrin 2014 → Sinai 2017 → this project 2026). Currently being updated by a parallel agent; see that file for the state-of-the-art treatment.

findings/the-perfect-flow-essay.md — the literary-critical counterpart. Argues that the Quran, read linearly, performs a *perfect flow* at the pericope scale even where whole-surah and whole-muṣḥaf ring claims fail. A humanistic-register complement to the quantitative monograph, rooted in *‘ilm al-balāgha* and *‘ilm al-munāsabāt*, extensively cross-referenced to al-Biqā'ī's *Nazm al-Durar*.

THE-MAN-AT-THE-CENTER.md — the historical-anthropological treatment. Asks: if the text has this compositional structure, what can we infer about the person-or-community at its center? Treats Muḥammad as a historical subject and reads the text's structural self-disclosures (the four *Muḥammad* occurrences all being Medinan post-Hijra; the chronological decline of Moses-mentions; the post-Hijra rhetorical shift from Lord-addressing to community-law) as evidence for an authorial trajectory independent of the *i'jāz* literature. The most speculative of the three volumes; framed as hypothesis, not conclusion.

7. What to read next

A skeptical reader with thirty minutes has read it. A reader with three hours should read **findings/scholar-commentary.md** (6,460-word narrative synthesis) and skim **docs/master-index.md**. A reader with thirty hours should read the monograph end-to-end and spot-check two or three phase-B findings against their own preferred null model. A reader with three hundred hours should replicate the project from **/data/** and **/analysis/tools/** — every rule is locked, every anchor value is test-covered, and the full audit trail is in **journal/**.

The project's central epistemic claim is modest. The Quran is not a cipher awaiting numerological decryption; it is a literarily dense composition whose coherence survives at the pericope scale, whose alleged global miracles mostly do not survive audit, and whose classical exegetical tradition was substantially correct about the *kinds* of structure present. The moral for the skeptic: most miracles collapse. The moral for the devotee: what survives is the composition itself — which is, after all, what the tradition has always pointed to.

End of Executive Summary. Approximately 3,050 words.